

Refining your plots, reports and code

MSc in Statistics 2025/26, Imperial College London

Zak Varty

Welcome Back!

Today's objectives

1. use R code within Quarto for single-document literate programming tasks
2. critique and refine default R outputs (such as plots and tables) for reports
3. understand how and when to show code within reports
4. use style guides and directory templates to provide a consistent structure across projects

Literate Programming

What is literate programming?

- A paradigm that combines:
 - **Code** (that executes)
 - **Narrative** (that explains)
- Result: documents that are both **human-readable** and **machine-executable**

Adding code blocks

```
```{r}
use :: to access data or functions from within a package
penguins <- palmerpenguins::penguins
```
```

```
# use :: to access data or functions from within a package
penguins <- palmerpenguins::penguins
```

Unwanted Warnings

Sometimes you get warnings that you don't need to draw attention to

```
library(palmerpenguins)
```

Attaching package: 'palmerpenguins'

The following objects are masked from 'package:datasets':

```
penguins, penguins_raw
```

Suppressing Warnings

```
```{r}
#| warning: false
library(palmerpenguins)
```
```

```
library(palmerpenguins)
```

Code blocks that show but don't run


Sometimes you want people to do as you say, not as you do. For example, I might want to ask you to install a package but without downloading it every time these notes are rendered.

```
```{r}
#| eval: false
install.packages("palmerpengins")
```
```

```
install.packages("palmerpengins")
```

Code blocks that run but don't show

```
```{r}
#| echo: false
plot(x = penguins$bill_length_mm,
 y = penguins$bill_depth_mm)
```
```



Multiple outputs

When you have multiple objects being printed in a single block, these will often get split up when the document is rendered.

```
```{r}
mean(penguins$body_mass_g, na.rm = TRUE)
median(penguins$body_mass_g, na.rm = TRUE)
```
```

```
mean(penguins$body_mass_g, na.rm = TRUE)
```

```
[1] 4201.754
```

```
median(penguins$body_mass_g, na.rm = TRUE)
```

```
[1] 4050
```

Multiple outputs (hold)

Top keep this displaying as a single block, you can set `results` to `hold`.

```
```{r}
#| results: hold
mean(penguins$body_mass_g, na.rm = TRUE)
median(penguins$body_mass_g, na.rm = TRUE)
```
```

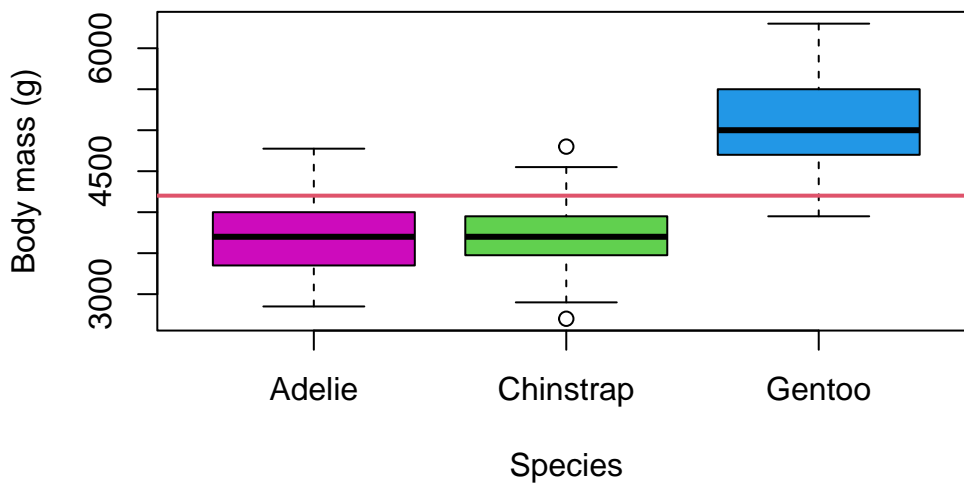
```
mean(penguins$body_mass_g, na.rm = TRUE)
median(penguins$body_mass_g, na.rm = TRUE)
```

```
[1] 4201.754
```

```
[1] 4050
```

Motivating Example

To give us an example to play with, we could consider the following question: are the average weights of all species of penguin the same?



Tables

How many penguins of each species do we have?

```
table(penguins$species)
```

```

      Adelie Chinstrap   Gentoo
      152       68      124

```

```

# fit linear model
mass_lm <- lm(formula = body_mass_g ~ 1 + species, penguins)

# extract coefficients and standard errors
mass_lm_coefficients <- summary(mass_lm)$coefficients[,1:2]
mass_lm_coefficients

```

```

              Estimate Std. Error
(Intercept)  3700.66225    37.61935
speciesChinstrap  32.42598    67.51168
speciesGentoo   1375.35401    56.14797

```

Nicer Tables

Example 1

```

species_counts <- table(penguins$species)
knitr::kable(species_counts, col.names = c("Species", "Count"))

```

| Species | Count |
|-----------|-------|
| Adelie | 152 |
| Chinstrap | 68 |
| Gentoo | 124 |

Example 2

```
knitr::kable(x = mass_lm_coefficients)
```

| | Estimate | Std. Error |
|------------------|------------|------------|
| (Intercept) | 3700.66225 | 37.61935 |
| speciesChinstrap | 32.42598 | 67.51168 |
| speciesGentoo | 1375.35401 | 56.14797 |

We could set the row names to improve formatting.

```
row.names(mass_lm_coefficients) <- c("Intercept", "Chinstrap", "Gentoo")
knitr::kable(x = mass_lm_coefficients)
```

| | Estimate | Std. Error |
|-----------|------------|------------|
| Intercept | 3700.66225 | 37.61935 |
| Chinstrap | 32.42598 | 67.51168 |
| Gentoo | 1375.35401 | 56.14797 |

And then refine the column names.

```
knitr::kable(
  x = mass_lm_coefficients,
  col.names = c("MLE", "Std. Err"))
```

| | MLE | Std. Err |
|-----------|------------|----------|
| Intercept | 3700.66225 | 37.61935 |
| Chinstrap | 32.42598 | 67.51168 |
| Gentoo | 1375.35401 | 56.14797 |

Finally, we really don't need to be showing all of those decimal places.

```
knitr::kable(
  x = round(mass_lm_coefficients, 2),
  col.names = c("MLE", "Std. Err"))
```


| | MLE | Std. Err |
|-----------|---------|----------|
| Intercept | 3700.66 | 37.62 |
| Chinstrap | 32.43 | 67.51 |
| Gentoo | 1375.35 | 56.15 |

A Semi-scripted approach

In report-style assessments you will not usually display your code in the main body of the report. In that case you can still use literate programming but you should:

- hide all of your code blocks
- clearly section your code with structural comments
- include an extra code block at the end of the report (this is in the template by default).

```
#####  
# Question 3  
#####  
  
# Part A -----  
  
# simulate data  
n <- 100  
x <- rnorm(n, mean = 0, sd = 1)  
y <- rnorm(n, mean = x, sd = abs(x) / 2)  
type <- ifelse(x > 0, "A", "B")  
  
# collect and clean workspace  
sim <- data.frame(x, y, type)  
rm(n, x, y, type)  
  
# Part B -----  
plot(x = sim$x, y = sim$y)
```



Labeling and Referencing

Referencing TLDR

- Label items of markdown using the `{#type-label}`
- Label code output with the `#| label: type-label`
- Reference either in the main text using `@type-label`

Example reference types: `sec`, `fig`, `tbl`, `eq`.

For full documentation see [Quarto Cross-References](#).

Referencing Sections and Subsections

```
# Labeling and Referencing {#sec-labeling}  
...  
@sec-labeling focuses on labelling and referencing parts of your document.
```

Section focuses on labelling and referencing parts of your document.

Referencing Tables


```

```{r}
#| label: tbl-palmer
#| tbl-cap: "Example rows of the Palmer penguin dataset."

library(knitr)
kable(head(penguins))
```

```

```

library(knitr)
kable(head(penguins)[ ,1:6])

```

Table 6: Example rows of the Palmer penguin dataset.

| species | island | bill_length_mm | bill_depth_mm | flipper_length_mm | body_mass_g |
|---------|-----------|----------------|---------------|-------------------|-------------|
| Adelie | Torgersen | 39.1 | 18.7 | 181 | 3750 |
| Adelie | Torgersen | 39.5 | 17.4 | 186 | 3800 |
| Adelie | Torgersen | 40.3 | 18.0 | 195 | 3250 |
| Adelie | Torgersen | NA | NA | NA | NA |
| Adelie | Torgersen | 36.7 | 19.3 | 193 | 3450 |
| Adelie | Torgersen | 39.3 | 20.6 | 190 | 3650 |

Table 6 gives example data entries from the Palmer penguins dataset. In your reports, all tables should have a caption and be discussed in the main text.

Referencing Figures

```

plot(x = sim$x, y = sim$y)

```




Figure 1

Simulating data can be an important tool to understand how models behave in ideal conditions. Figure 1 shows an example of some bivariate data, simulated in Section .

Referencing Equations

```
$$ e^{\pi i} + 1 = 0. $$#eq-euler
```

$$e^{\pi i} + 1 = 0. \tag{1}$$

Equation 1 states an elegant and well known result. By convention, equations should only be numbered if they are mentioned in the main text.

Citing Sources

Add a bibliography to your document using the bibliography YAML metadata field. As for LaTeX this should point to a `.bib` file, detailing the sources you will cite. I have provided a small example bibliography file `ref.bib`.

```
---
title: "Refining your plots, reports and code"
bibliography: ref.bib
---
```

By default, Pandoc will automatically generate a list of works cited and place it in the document if the style calls for it. It will be placed in a div with the id refs if one exists:

```
### References

::: {#refs}
:::
```

To cite books or articles you can use a direct citation to talk about a source such as Box and Cox (1964) within a sentence, or a parenthetical reference to evidence a more general claim (Casella and Berger 2024) .

Markdown Syntax:


- @box1964analysis gives Box and Cox (1964);
- [casella2024statistical] gives (Casella and Berger 2024).

Further information in the [citation documentation](#).

Refining Your Figures


A basic plot

```
plot(
  x = penguins$bill_length_mm,
  y = penguins$bill_depth_mm)
```



Titles and Subtitles


```
plot(  
  x = penguins$bill_length_mm,  
  y = penguins$bill_depth_mm,  
  main = "Penguin bill dimensions",  
  sub = "Data source: Palmer Penguins")
```



Axis Labels


```
plot(  
  x = penguins$bill_length_mm,  
  y = penguins$bill_depth_mm,  
  main = "Penguin bill dimensions",  
  xlab = "Length (mm)",  
  ylab = "Depth (mm)")
```

Penguin bill dimensions



Resizing Axis Labels

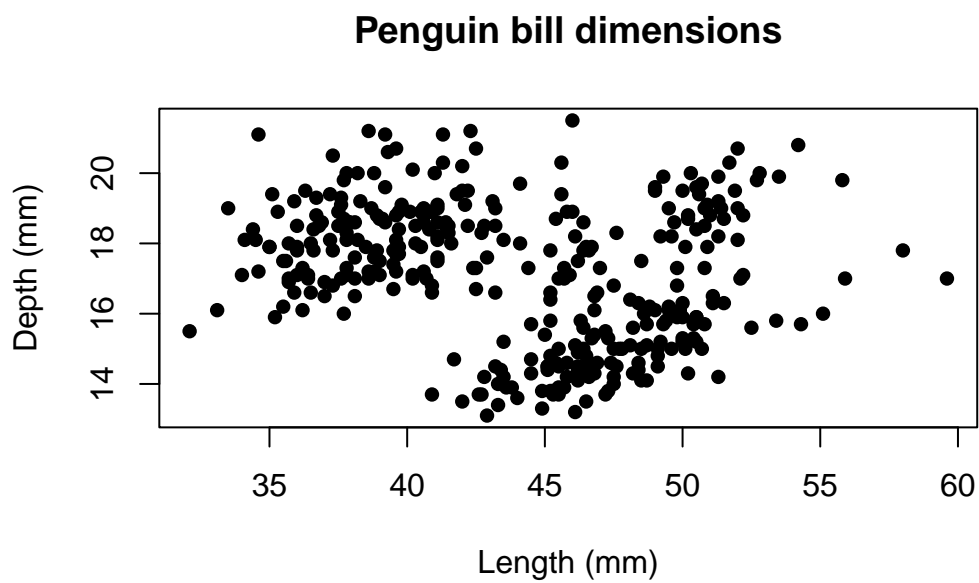
```
plot(  
  x = penguins$bill_length_mm,  
  y = penguins$bill_depth_mm,  
  xlab = "Length (mm)",  
  ylab = "Depth (mm)",  
  cex.axis = 1.2, # relative scaling of axis values  
  cex.lab = 1.4) # relative scaling of axis labels
```



Visual Grammar


Plotting Characters

```
plot(  
  x = penguins$bill_length_mm,  
  y = penguins$bill_depth_mm,  
  main = "Penguin bill dimensions",  
  xlab = "Length (mm)",  
  ylab = "Depth (mm)",  
  pch = 16)
```



```
plot(  
  x = penguins$bill_length_mm,  
  y = penguins$bill_depth_mm,  
  main = "Penguin bill dimensions",  
  xlab = "Length (mm)",  
  ylab = "Depth (mm)",  
  pch = 16,  
  col = "blue")
```


Penguin bill dimensions



Over-plotting happens when one point or line in a figure obscures another. We can resolve this issue by picking a transparent colour using `rgb()`.

```
plot(  
  x = penguins$bill_length_mm,  
  y = penguins$bill_depth_mm,  
  main = "Penguin bill dimensions",  
  xlab = "Length (mm)",  
  ylab = "Depth (mm)",  
  pch = 16,  
  col = rgb(0,0,0,0.5))
```

Penguin bill dimensions



We could also use colour to encode another aspect of our data visually.

```
plot(  
  x = penguins$bill_length_mm,  
  y = penguins$bill_depth_mm,  
  main = "Penguin bill dimensions",  
  xlab = "Length (mm)",  
  ylab = "Depth (mm)",  
  pch = 16,  
  col = penguins$species)
```

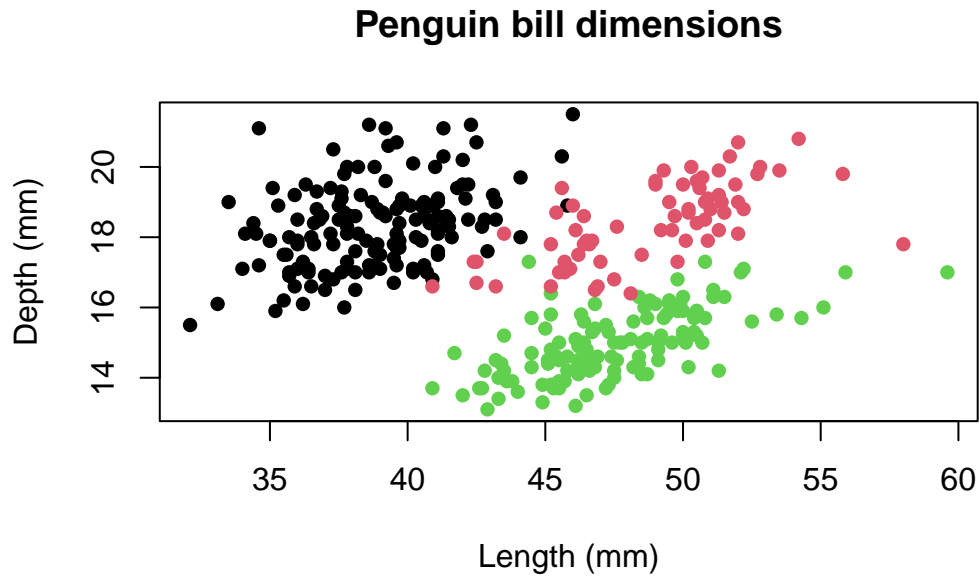



Figure 2: Penguin bill dimensions by species: Adelie (black), Chinstrap (red) and Gentoo (green).

To make the plot more accessible, we could also alter the plotting character for each species.

```
plot(  
  x = penguins$bill_length_mm,  
  y = penguins$bill_depth_mm,  
  main = "Penguin bill dimensions",  
  xlab = "Length (mm)",  
  ylab = "Depth (mm)",  
  pch = 15 + as.numeric(penguins$species),  
  col = penguins$species)
```




Figure 3: Penguin bill dimensions by species: Adelie (black circles), Chinstrap (red triangles) and Gentoo (green diamonds).

If we are plotting lines instead, we can also change the line width and type.

```
time <- seq(from = 0, to = 100, by = 0.01)
displacement <- 10 + 5 * sin(time)
```

```
plot(
  x = time,
  y = displacement,
  type = "l",
  bty = "n",
  col = "blue",
  lwd = 2, # line width
  lty = 2) # dashed line
```





Figure 4: Displacement of an undamped simple harmonic oscillator.

Layering your plots


We can overlay a line on our previous plot using `lines()`.

```
plot(  
  x = time,  
  y = displacement,  
  type = "l",  
  bty = "n",  
  col = "blue",  
  lwd = 2, # line width  
  lty = 2) # dashed line  
spring_2 = 10 + 3 * sin(0.25 * time)  
lines(x = time, y = spring_2, lwd = 2)
```



Similarly, we can add one or more points to the plot by using `points()`.

```
plot(
  x = time,
  y = displacement,
  type = "l",
  bty = "n",
  col = "blue",
  lwd = 2,
  lty = 2)
lines(x = time, y = spring_2, lwd = 2)
points(x = c(2, 10, 18, 26) * pi, y = rep(13, 4), pch = 16, col = "red")
```



Finally, we might want to add a legend layer to our earlier scatterplot in Figure 3.

```

plot(
  x = penguins$bill_length_mm,
  y = penguins$bill_depth_mm,
  col = penguins$species,
  pch = 15 + as.numeric(penguins$species),
  xlim = c(29,70), # extend plotting area to make room for legend
  ylim = c(12,23),
  main = "Penguin bill dimensions",
  xlab = "Length (mm)",
  ylab = "Depth (mm)",
  bty = "n")
legend(
  "bottomright",
  title = "Species",
  legend = c("Adelie", "Chinstrap", "Gentoo"),
  pch = 16:18,
  col = 1:3,
  bty = "n")

```




Figure 5: Penguin bill dimensions by species.

Combining your plots

```
```{r}
#| layout-ncol: 2
#| label: fig-example-plots
#| fig-cap: "Two example plots relating to the Gaussian distribution."
#| fig-subcap:
#| - "Gaussian distribution"
#| - "Gaussian QQ-plot"

x <- seq(from = -3, to = 3, by = 0.01)
density <- dnorm(x, mean = 0, sd = 1)

plot(
 x = x,
 y = density,
 bty = "n",
 cex.axis = 1.4,
 cex.lab = 1.4,
 type = "l",
 lwd = 2)

y <- rnorm(n = 100, mean = 3, sd = 2)
qqnorm(y, cex.axis = 1.4, cex.lab = 1.4, main = "")
```
```

```
x <- seq(from = -3, to = 3, by = 0.01)
density <- dnorm(x, mean = 0, sd = 1)

plot(
  x = x,
  y = density,
  bty = "n",
  cex.axis = 1.4,
  cex.lab = 1.4,
  type = "l",
  lwd = 2)

y <- rnorm(n = 100, mean = 3, sd = 2)
qqnorm(y, cex.axis = 1.4, cex.lab = 1.4, main = "")
```

Figure 6 shows how we can create multiple plots within a code block and then have these appear




Figure 6: Two example plots relating to the Gaussian distribution.

within the same figure as sub-figures. A probability density function is shown in Figure 6a, while a quantile-quantile plot is shown in Figure 6b.

What goes where? Tiles, Captions and Descriptions

| Aspect | Plot Title | Figure Caption |
|-------------|-------------------------------|--|
| Location | Inside or above the plot | Outside the plot |
| Purpose | Quick context for the plot | Formal description with interpretive context |
| Referencing | Rarely referenced in the text | Commonly referenced (e.g., “see Figure 1”) |

Where a figure or table is mentioned in the main text you should explain *how* it supports any claims that you are making.

BAD: Figure 3 demonstrates Simpson’s paradox.

GOOD: Figure 3 shows that while there appears to be a positive association between bill length and depth *within* each species, this association disappears or perhaps reverses when considering all species together. This is an example of Simpson’s paradox [simpson1951contingency].

Useful Resources

- Imperial Coursework Report Template
- [Quarto documentation](#)
- [Make a reprex, please](#)
- Telling Stories With Data
 - [3 - Reproducible workflows](#)
- R style guides [Google](#), [Tidyverse](#).
- Effective Data Science
 - [1 - Organising your work](#)
 - [8 - Exploratory data analysis](#)
 - [9 - Data visualisation](#)

- [The TidyTuesday Cookbook](#)

Review

1. use R code within Quarto for single-document literate programming tasks
2. critique and refine default R outputs (such as plots and tables) for reports
3. understand how and when to show code within reports
4. use style guides and directory templates to provide a consistent structure across projects

By default, Pandoc will automatically generate a list of works cited and place it in the document if the style calls for it. It will be placed in a div with the id `refs` if one exists:

References

Box, George EP, and David R Cox. 1964. “An Analysis of Transformations.” *Journal of the Royal Statistical Society Series B: Statistical Methodology* 26 (2): 211–43.

Casella, George, and Roger Berger. 2024. *Statistical Inference*. Chapman; Hall/CRC.